

Ants and Bees - Classification to Demonstrate Transfer Learning

- This notebook is supposed to be an amateur's first foray into Transfer Learning.
- It operates on a dataset that is relatively small - with 120 images for ants and bees. There are 75 validation images for each class as well.
- Usually, this is a very small dataset to generalize upon, if trained from scratch. Since we are using transfer learning, we should be able to generalize reasonably well.

Downloading the Data

In []:

```
from io import BytesIO
from typing import Union
from zipfile import ZipFile
import requests
from os import path, rename

"""Downloads a zip file from `url`, extracts it at `dest`, renames the extracted folder if `new_dir_name` is provided"""

def download_zip_and_extract(url: str, dest: str, new_dir_name: Union[str, None] = None):
    file_name = ".".join(url.split('/')[-1].split('.')[0:-1])
    dest_path = path.join(path.curdir, dest)

    if path.exists(path.join(dest_path, new_dir_name if new_dir_name else file_name)):
        print("Destination folder already exists")
        return

    print("Downloading started")
    req=requests.get(url)
    print("Downloading completed. Extraction started")

    zipfile=ZipFile(BytesIO(req.content))
    zipfile.extractall(dest_path)

    print("Extraction completed")

    if new_dir_name:
        rename(path.join(dest_path, file_name), new_dir_name)
```

In []:

```
DATA_URL = "https://download.pytorch.org/tutorial/hymenoptera_data.zip"
data_dest = "."
download_zip_and_extract(DATA_URL, data_dest, "data")
```

Destination folder already exists

Relevant Imports

```
In [ ]: import torch
from torch import nn, optim
from torch.optim import lr_scheduler
from torch.utils.data import DataLoader
import numpy as np
import torchvision
from torchvision import datasets, transforms as T, models
import matplotlib.pyplot as plt
import time
import os

plt.ion()
```

```
Out[ ]: <matplotlib.pyplot._IonContext at 0x20f896af820>
```

Load Data

```
In [ ]: data_transforms = {
    'train': T.Compose([
        T.RandomResizedCrop(224),
        T.RandomHorizontalFlip(),
        T.ToTensor(),
        T.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
    ]),
    'val': T.Compose([
        T.Resize(256),
        T.CenterCrop(224),
        T.ToTensor(),
        T.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
    ])
}
```

```
In [ ]: batch_size = 4
```

```
In [ ]: data_dir = path.join(data_dest, "data")

image_datasets = {x: datasets.ImageFolder(
    os.path.join(data_dir, x), transform=data_transforms[x]) for x in ['train', 'val']}
dataloaders = {x: DataLoader(image_datasets[x], batch_size=batch_size, shuffle=(
    True if x == 'train' else False), num_workers=4) for x in ['train', 'val']}

dataset_sizes = {x: len(image_datasets[x]) for x in ['train', 'val']}
category_names = image_datasets['train'].classes
```

```
In [ ]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
device
```

```
Out[ ]: device(type='cuda')
```

Visualising a Few Images

- We'll plot a few images to see the effects of the data transforms.

In []:

```
from torch import Tensor

def imshow(inp: Tensor, title=None):
    """Plot image if `inp` is a tensor"""

    inp = inp.cpu().numpy().transpose((1, 2, 0))
    mean = np.array([0.485, 0.456, 0.406])
    std = np.array([0.229, 0.224, 0.225])

    inp = std * inp + mean
    np.clip(inp, 0, 1)

    if title is not None:
        plt.title(title)

    return plt.imshow(inp)
```

In []:

```
# Get a batch of training data
inputs, classes = next(iter(dataloaders['train']))

# make a grid from batch
out = torchvision.utils.make_grid(inputs)

imshow(out, [category_names[x] for x in classes])
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Out[]:



Training the Model

It will help us to:

- schedule the learning rate
- saving the best model

In []:

```
from copy import deepcopy
from torch.nn import Module

def train_model(model: Module, criterion, optimizer: optim.Optimizer, scheduler, num_
since = time.time()

best_model_wts = deepcopy(model.state_dict())
best_acc = 0.0

for epoch in range(num_epochs):
    print(f"Epoch: {epoch} / {num_epochs}")

    for phase in ['train', 'val']:
        if phase == 'train':
            model.train()

        else:
            model.eval()

        running_loss, running_corrects = 0.0, 0

        for inputs, labels in dataloaders[phase]:
            inputs, labels = inputs.to(device), labels.to(device)

            # zero all the layers' parameters gradients
            optimizer.zero_grad()

            # feed forward
            with torch.set_grad_enabled(phase == 'train'):
                outputs = model(inputs)
                _, preds = torch.max(outputs, dim=1)
                loss = criterion(outputs, labels)

            # backpropagate loss only if training
            if phase == 'train':
                loss.backward()
                optimizer.step()

            # default loss item is mean, therefore we multiply it with the number of items
            running_loss += loss.item() * inputs.size(0)
            running_corrects += torch.sum(preds == labels)

        if phase == 'train':
            scheduler.step()

        epoch_loss = running_loss / dataset_sizes[phase]
        epoch_acc = running_corrects.double() / dataset_sizes[phase]

        print(f'{phase} Loss: {epoch_loss:.4f} Acc: {epoch_acc:.4f}')

        # deep copy the model
        if phase == 'val' and epoch_acc > best_acc:
            best_acc = epoch_acc
            best_model_wts = deepcopy(model.state_dict())

print()
```

```

time_elapsed = time.time() - since
print(
    f'Training complete in {time_elapsed // 60:.0f}m {time_elapsed % 60:.0f}s')
print(f'Best val Acc: {best_acc:.4f}')

model.load_state_dict(best_model_wts)
return model

```

Visualising Model Predictions

In []:

```
from matplotlib.image import AxesImage
```

```

def process_tensor_to_display(inp: Tensor) -> np.ndarray:
    inp = inp.cpu().numpy().transpose((1, 2, 0))
    mean = np.array([0.485, 0.456, 0.406])
    std = np.array([0.229, 0.224, 0.225])

    inp = std * inp + mean
    np.clip(inp, 0, 1)

    return inp

```

In []:

```

def visualise_predictions(model: Module, num_images: int, n_rows: int, n_cols: int):
    was_training = model.training
    model.eval()
    images_so_far = 0
    fig, ax = plt.subplots(n_rows, n_cols)

    with torch.no_grad():
        for i, (inputs, labels) in enumerate(dataloaders['val']):
            inputs, labels = inputs.to(device), labels.to(device)

            outputs = model(inputs)
            _, preds = torch.max(outputs, dim=1)

            for j in range(inputs.size(0)):

                # below call takes number of rows, number of cols, index of item that starts
                plotX, plotY = images_so_far // n_cols, images_so_far % n_cols
                ax[plotX, plotY].imshow(process_tensor_to_display(inputs[j].detach()))
                ax[plotX, plotY].axis('off')
                ax[plotX, plotY].set_title(f'{category_names[preds[j]]}')
                images_so_far += 1

                # imshow(inputs[j].detach())

            if images_so_far == num_images:
                model.train(was_training)
                return

    model.train(was_training)

```

Finetuning the convnet

Load a pretrained model and reset final fully connected layer

```
In [ ]: model_ft = models.resnet18(pretrained=True)
num_ftrs = model_ft.fc.in_features

model_ft.fc = nn.Linear(num_ftrs, len(category_names))

model_ft = model_ft.to(device)
```

```
In [ ]: criterion = nn.CrossEntropyLoss()

# Observe that all parameters are being optimized
optimizer_ft = optim.SGD(model_ft.parameters(), lr=0.001, momentum=0.9)

# Decay LR by a factor of 0.1 every 7 epochs
exp_lr_scheduler = lr_scheduler.StepLR(optimizer_ft, 7, gamma=0.1)
```

```
In [ ]: model_ft = train_model(model_ft, criterion, optimizer_ft, exp_lr_scheduler, num_epochs)
```

```
Epoch: 0 / 30
d:\ProgramFiles\anaconda\envs\fastapi-ml\lib\site-packages\torch\nn\functional.py:71
8: UserWarning: Named tensors and all their associated APIs are an experimental feature and subject to change. Please do not use them for anything important until they are released as stable. (Triggered internally at ..\c10\core/TensorImpl.h:1156.)
    return torch.max_pool2d(input, kernel_size, stride, padding, dilation, ceil_mode)
train Loss: 0.5452 Acc: 0.6844

val Loss: 0.2099 Acc: 0.9216

Epoch: 1 / 30
train Loss: 0.4097 Acc: 0.8279

val Loss: 0.2839 Acc: 0.8824

Epoch: 2 / 30
train Loss: 0.5586 Acc: 0.7746

val Loss: 0.3774 Acc: 0.8366

Epoch: 3 / 30
train Loss: 0.5762 Acc: 0.7418

val Loss: 0.2786 Acc: 0.8889

Epoch: 4 / 30
train Loss: 0.4245 Acc: 0.8033

val Loss: 0.3033 Acc: 0.8889

Epoch: 5 / 30
train Loss: 0.5666 Acc: 0.7746

val Loss: 0.3868 Acc: 0.8497
```

```
Epoch: 6 / 30
train Loss: 0.3811 Acc: 0.8648

val Loss: 0.3813 Acc: 0.8889

Epoch: 7 / 30
train Loss: 0.4489 Acc: 0.7951

val Loss: 0.2856 Acc: 0.9150

Epoch: 8 / 30
train Loss: 0.3303 Acc: 0.8525

val Loss: 0.2610 Acc: 0.9281

Epoch: 9 / 30
train Loss: 0.4298 Acc: 0.8361

val Loss: 0.2819 Acc: 0.9020

Epoch: 10 / 30
train Loss: 0.3481 Acc: 0.8320

val Loss: 0.2554 Acc: 0.9216

Epoch: 11 / 30
train Loss: 0.2995 Acc: 0.8852

val Loss: 0.2408 Acc: 0.9281

Epoch: 12 / 30
train Loss: 0.3145 Acc: 0.8566

val Loss: 0.2455 Acc: 0.9216

Epoch: 13 / 30
train Loss: 0.3038 Acc: 0.8361

val Loss: 0.2469 Acc: 0.9216

Epoch: 14 / 30
train Loss: 0.2442 Acc: 0.8975

val Loss: 0.2867 Acc: 0.8954

Epoch: 15 / 30
train Loss: 0.2736 Acc: 0.8648

val Loss: 0.2582 Acc: 0.9020

Epoch: 16 / 30
train Loss: 0.3095 Acc: 0.8525

val Loss: 0.2468 Acc: 0.9216

Epoch: 17 / 30
train Loss: 0.2819 Acc: 0.8730

val Loss: 0.2432 Acc: 0.9281
```

```
Epoch: 18 / 30
train Loss: 0.2765 Acc: 0.8607

val Loss: 0.2443 Acc: 0.9216

Epoch: 19 / 30
train Loss: 0.2397 Acc: 0.8852

val Loss: 0.2364 Acc: 0.9281

Epoch: 20 / 30
train Loss: 0.3606 Acc: 0.8279

val Loss: 0.2317 Acc: 0.9346

Epoch: 21 / 30
train Loss: 0.3057 Acc: 0.8770

val Loss: 0.2963 Acc: 0.9216

Epoch: 22 / 30
train Loss: 0.2689 Acc: 0.8811

val Loss: 0.2381 Acc: 0.9216

Epoch: 23 / 30
train Loss: 0.2571 Acc: 0.9016

val Loss: 0.2515 Acc: 0.9216

Epoch: 24 / 30
train Loss: 0.2707 Acc: 0.8934

val Loss: 0.2342 Acc: 0.9216

Epoch: 25 / 30
train Loss: 0.3202 Acc: 0.8648

val Loss: 0.2300 Acc: 0.9281

Epoch: 26 / 30
train Loss: 0.2434 Acc: 0.9057

val Loss: 0.2707 Acc: 0.9150

Epoch: 27 / 30
train Loss: 0.3011 Acc: 0.8689

val Loss: 0.2687 Acc: 0.9020

Epoch: 28 / 30
train Loss: 0.2325 Acc: 0.9098

val Loss: 0.2501 Acc: 0.9216

Epoch: 29 / 30
train Loss: 0.2538 Acc: 0.9057
```

```
val Loss: 0.2400 Acc: 0.9281
```

Training complete in 2m 34s

```
In [ ]: visualise_predictions(model_ft, 6, 2, 3)
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Predicted ants



Predicted bees



Predicted ants



Predicted ants



Predicted ants



Predicted ants



ConvNet as a Fixed Feature Extractor

- Here we will freeze all the layers of the network, except for the final layer.
- We need to set `requires_grad = False` for all those layers, so that their loss is not computed during the backward pass.

```
In [ ]: model_conv = models.resnet34(pretrained=True)

for param in model_conv.parameters():
    param.requires_grad = False
```

Downloading: "https://download.pytorch.org/models/resnet34-b627a593.pth" to C:\Users\janma\.cache\torch\hub\checkpoints\resnet34-b627a593.pth
100% [██████████] 83.3M/83.3M [02:23<00:00, 610kB/s]

```
In [ ]: num_ftrs = model_conv.fc.in_features
model_conv.fc = nn.Linear(num_ftrs, len(category_names))

model_conv.to(device)
```

```
Out[ ]: ResNet(
  (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
```

```
ue)
    (relu): ReLU(inplace=True)
    (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
e)
    (layer1): Sequential(
        (0): BasicBlock(
            (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
            (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
            (relu): ReLU(inplace=True)
            (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
            (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        )
        (1): BasicBlock(
            (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
            (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
            (relu): ReLU(inplace=True)
            (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
            (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        )
        (2): BasicBlock(
            (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
            (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
            (relu): ReLU(inplace=True)
            (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
            (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        )
    )
    (layer2): Sequential(
        (0): BasicBlock(
            (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
            (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
            (relu): ReLU(inplace=True)
            (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
            (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        )
        (downsample): Sequential(
            (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
            (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        )
    )
    (1): BasicBlock(
        (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
```

```
(bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_st
ts=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bi
as=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_st
ts=True)
    )
  (2): BasicBlock(
    (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bi
as=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_st
ts=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bi
as=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_st
ts=True)
    )
  (3): BasicBlock(
    (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bi
as=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_st
ts=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bi
as=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_st
ts=True)
    )
)
(layer3): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bi
as=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_st
ts=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bi
as=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_st
ts=True)
    (downsample): Sequential(
      (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_st
ts=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bi
as=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_st
ts=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bi
as=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_st
ts=True)
    )
)
```

```
(2): BasicBlock(
    (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
(3): BasicBlock(
    (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
(4): BasicBlock(
    (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
(5): BasicBlock(
    (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
)
(layer4): Sequential(
    (0): BasicBlock(
        (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
        (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
    (downsample): Sequential(
        (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
        (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
)
```

```
ts=True)
    )
)
(1): BasicBlock(
    (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
(2): BasicBlock(
    (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
)
(avgpool): AdaptiveAvgPool2d(output_size=(1, 1))
(fc): Linear(in_features=512, out_features=2, bias=True)
```

```
In [ ]: criterion = nn.CrossEntropyLoss()

optimizer = optim.SGD(model_conv.parameters(), 0.001, 0.9)

exp_lr_scheduler = lr_scheduler.StepLR(optimizer, 7, 0.01)
```

Train and Evaluate

```
In [ ]: model_conv = train_model(model_conv, criterion, optimizer, exp_lr_scheduler)
```

```
Epoch: 0 / 25
d:\ProgramFiles\anaconda\envs\fastapi-ml\lib\site-packages\torch\nn\functional.py:71
8: UserWarning: Named tensors and all their associated APIs are an experimental feature and subject to change. Please do not use them for anything important until they are released as stable. (Triggered internally at ..\c10\core/TensorImpl.h:1156.)
    return torch.max_pool2d(input, kernel_size, stride, padding, dilation, ceil_mode)
train Loss: 0.6462 Acc: 0.6598
```

```
val Loss: 0.2131 Acc: 0.9412
```

```
Epoch: 1 / 25
train Loss: 0.4524 Acc: 0.7828
```

```
val Loss: 0.3041 Acc: 0.8693
```

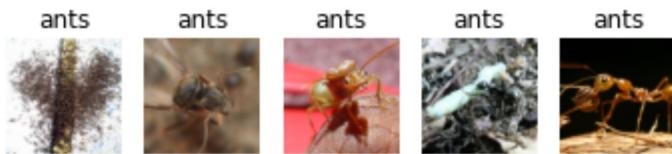
```
Epoch: 2 / 25
train Loss: 0.4432 Acc: 0.7992
```

```
val Loss: 0.1619 Acc: 0.9216  
Epoch: 3 / 25  
train Loss: 0.3744 Acc: 0.8279  
  
val Loss: 0.2635 Acc: 0.8824  
  
Epoch: 4 / 25  
train Loss: 0.4363 Acc: 0.8156  
  
val Loss: 0.2326 Acc: 0.9085  
  
Epoch: 5 / 25  
train Loss: 0.4172 Acc: 0.8156  
  
val Loss: 0.1654 Acc: 0.9281  
  
Epoch: 6 / 25  
train Loss: 0.2964 Acc: 0.8402  
  
val Loss: 0.2046 Acc: 0.9150  
  
Epoch: 7 / 25  
train Loss: 0.3512 Acc: 0.8525  
  
val Loss: 0.1634 Acc: 0.9346  
  
Epoch: 8 / 25  
train Loss: 0.2824 Acc: 0.8689  
  
val Loss: 0.1702 Acc: 0.9281  
  
Epoch: 9 / 25  
train Loss: 0.2675 Acc: 0.8770  
  
val Loss: 0.1388 Acc: 0.9412  
  
Epoch: 10 / 25  
train Loss: 0.3115 Acc: 0.8566  
  
val Loss: 0.1635 Acc: 0.9412  
  
Epoch: 11 / 25  
train Loss: 0.3183 Acc: 0.8607  
  
val Loss: 0.1524 Acc: 0.9412  
  
Epoch: 12 / 25  
train Loss: 0.2740 Acc: 0.8770  
  
val Loss: 0.1470 Acc: 0.9216  
  
Epoch: 13 / 25  
train Loss: 0.2642 Acc: 0.8689  
  
val Loss: 0.1840 Acc: 0.9281  
  
Epoch: 14 / 25
```

```
train Loss: 0.3747 Acc: 0.8566
val Loss: 0.1483 Acc: 0.9346
Epoch: 15 / 25
train Loss: 0.3237 Acc: 0.8648
val Loss: 0.1414 Acc: 0.9412
Epoch: 16 / 25
train Loss: 0.3138 Acc: 0.8730
val Loss: 0.1485 Acc: 0.9542
Epoch: 17 / 25
train Loss: 0.2813 Acc: 0.8811
val Loss: 0.1400 Acc: 0.9412
Epoch: 18 / 25
train Loss: 0.2980 Acc: 0.8648
val Loss: 0.1458 Acc: 0.9542
Epoch: 19 / 25
train Loss: 0.2572 Acc: 0.8975
val Loss: 0.1490 Acc: 0.9346
Epoch: 20 / 25
train Loss: 0.3071 Acc: 0.8648
val Loss: 0.2091 Acc: 0.9085
Epoch: 21 / 25
train Loss: 0.3139 Acc: 0.8689
val Loss: 0.1493 Acc: 0.9412
Epoch: 22 / 25
train Loss: 0.2518 Acc: 0.9016
val Loss: 0.1556 Acc: 0.9412
Epoch: 23 / 25
train Loss: 0.2769 Acc: 0.9016
val Loss: 0.1498 Acc: 0.9412
Epoch: 24 / 25
train Loss: 0.3111 Acc: 0.8566
val Loss: 0.1468 Acc: 0.9542
Training complete in 3m 37s
Rest val Acc: 0.954248
```

```
In [ ]: visualise_predictions(model_conv, 10, 2, 5)
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



In []: